



ENTERPRISE BUSINESS CASE

Agentic SDLC Factory

Value model, ROI and operating options for enterprise-scale delivery

\$386M net annual saving · 9.2x year-one ROI · payback in 1.2 months

Project start in minutes · one security plane across every system of record

Agentic SDLC Factory returns \$386M a year on a \$42M year-one investment

Situation - 6,000 projects a year, 25,000 engineers, \$2.81B of delivery labour.

Complication - AI assistance alone captures 12% of addressable effort; security and cost stay fragmented.

Resolution - one governed factory: 14 agents, 9 evidence gates, one security and cost plane.

\$386M

net annual saving, base case

9.2x

year-one ROI, payback 1.2 months

minutes

to start a project, versus 2-6 weeks

2.3 weeks

of delivery time removed per project

Source: illustrative model built from the assumptions on the next slide.

Four value claims, in order of importance

1

Speed to start a project

One requirements document in; provisioned project, repo and backlog out.

Minutes, versus 2-6 weeks

2

Speed of release

One ordered graph plans, builds, tests, secures and deploys.

2.3 weeks removed per project

3

Centralized security

Guardrails, allow-lists and token limits enforced server-side everywhere.

One plane, every system of record

4

Cost visibility

Every model call logs tokens to a project and an owner.

\$2,980 platform cost per project

Value scales with the portfolio, not with the team

5

Portfolio cost avoidance

\$298M-\$550M net a year across 6,000 projects, depending on the autonomy mode.

6

Consistency and reuse

The same 14 agents, 9 gates and model policy run on every project.

7

Audit-ready traceability

Append-only audit trail; work items carry commit, PR and deployment links.

8

Quality and improvement


Security stage before deploy; the Insights agent feeds Planning a ranked backlog.

Six assumptions drive the entire estimate

Enterprise scale	\$1T market cap; 25,000 FTE across engineering, QA, DevOps and delivery management
Cost base	\$190,000 fully loaded per FTE = \$4.75B total engineering base
Project portfolio	6,000 funded projects a year; 8 FTE for 16 weeks = \$467,692 per project = \$2.81B a year
Addressable effort	45% of project effort - requirements, planning, design, code, tests, security prep, release and ops docs = \$210,462 per project
Platform run cost	\$17.9M a year = \$2,980 per project, of which \$450 is model tokens
Year-1 investment	\$42M = run cost + \$18M enablement and change management + \$6M build, connectors and hardening

Illustrative planning model, not a quotation. Validate the assumptions with a 10-project pilot.

Net annual saving rises from \$134M to \$550M as gates move from people to policy



Manual + AI \$134M • Human Review \$298M • Minimal Human Review \$386M • Fully Autonomous \$550M net of \$17.9M platform run cost, across 6,000 projects a year

Minimal Human Review is the default: four gates, \$386M, 9.2x return

Manual + AI today	Human Review 9 of 9 gates	Minimal Review 4 of 9 gates · DEFAULT	Autonomous 0 of 9 gates
Gates ad hoc Uplift 12% Per project \$22,275	Gates 9 of 9 Uplift 25% Per project \$49,635	Gates 4 of 9 Uplift 32% Per project \$64,368	Gates none Uplift 45% Per project \$91,728
\$134M	\$298M	\$386M	\$550M
3.2x ROI · payback 3.3 mo	7.1x ROI · payback 1.6 mo	9.2x ROI · payback 1.2 mo	13.1x ROI · payback 0.9 mo
Baseline today	Regulated, high-risk work	Default enterprise balance	Trusted repeatable patterns

Net annual saving across 6,000 projects a year, after the \$17.9M platform run cost.

Five capabilities carry the platform

1

14 specialized Foundry agents

One deterministic graph, ordered 010 to 120, on a shared context bus.

2

Nine evidence-aware approval gates

Three autonomy policies per project; evidence checked at every gate.

3

Every model call through Azure API Management

Managed identity to Foundry; no browser or API bypass; every call logged.

4

Governed connectors

Agents never own data; each write lands in a system of record with an evidence URL.

5

Model choice and engine switching

Per-agent model dropdown, Model Router, and Foundry agents or GitHub Copilot.

Five more capabilities that de-risk the rollout

6

Human-in-the-loop editing

Reviewers select, edit, add or remove items before publication.

7

Durable execution

Persisted queue, retries, circuit breaker and restart recovery.

8

Parallel project sets

One wizard launches 2-20 projects with isolated runs and audit.

9

Iterative revisions

Restart from Requirements, Planning or Architecture, carrying approved gates forward.

10

Safe lifecycle and cost governance

Cascade cleanup of created resources; token limits and role permissions server-side.

Four systems of record, one security plane

AZURE DEVOPS
Work items, test plans,
pipelines, wiki
LIVE



GITHUB
Repos, branches, pull requests,
Actions
LIVE



JIRA
Work-item system of record
EXTENSIBLE



BITBUCKET
Source control system of record
EXTENSIBLE



CENTRALIZED SECURITY - ENFORCED IN ONE PLACE

Entra ID and RBAC · APIM gateway with managed identity · guardrails and tool allow-lists · token limits · append-only audit - server-side, identically, for every connector above

COST VISIBILITY - TOKENS BY PROJECT AND BY OWNER

Every call logs prompt, agent, model, token estimate, project, user, gate, run and correlation ID, so spend rolls up to a named project owner.

Token spend is attributable per project owner

\$2,980

platform cost per project

\$450

of that is model tokens

9

identifiers logged on every model call

Platform cost line	Annual	Per project
Model tokens	\$2.7M	\$450
Infrastructure - APIM, App Service, Cosmos	\$1.2M	\$200
Platform team - 12 FTE	\$2.3M	\$380
Coding-assistant seats - 25,000	\$11.7M	\$1,950
Total run cost	\$17.9M	\$2,980

Grounding the agents in company standards is the next increment

1

Azure API Center

Generated code calls real, versioned, governed contracts instead of invented endpoints.

2

Data catalog and dictionary

Data models and queries follow approved entities, naming standards and schemas.

3

Brand and design-system site


Generated front ends use approved components, tokens, typography and accessibility rules.

4

Additional systems of record

Jira, Bitbucket and ITSM inherit the same security, audit and cost telemetry plane.

Every project runs the same nine-gate graph and loops back to Planning



NINE EVIDENCE GATES

Plan and Scope · Backlog Generation · Architecture and Design · Code Generation · Pull Request Review · Test Acceptance · Security Review · Release and Deployment · Operate and Improve



The Insights agent returns measurable thresholds and a ranked backlog to Planning. A completed project can restart a traceable revision from Requirements, Planning or Architecture.

THE ASK

Fund a 10-project pilot in Minimal Human Review

- 1 Validate the 45% addressable share and the 32% uplift on real projects.
- 2 Prove the security plane and per-owner cost rollup on Azure DevOps and GitHub first.
- 3 Set autonomy per project class: Human Review, Minimal as default, Autonomous.

\$386M a year · 9.2x year-one ROI · payback in 1.2 months

Financial figures are an illustrative planning model, not a quotation.